

于延辰


InnoDB加锁分析

Background

- ❖ MVCC
- ❖ Clustered and Secondary Indexes
- ❖ 2PL
- ❖ Isolation Level

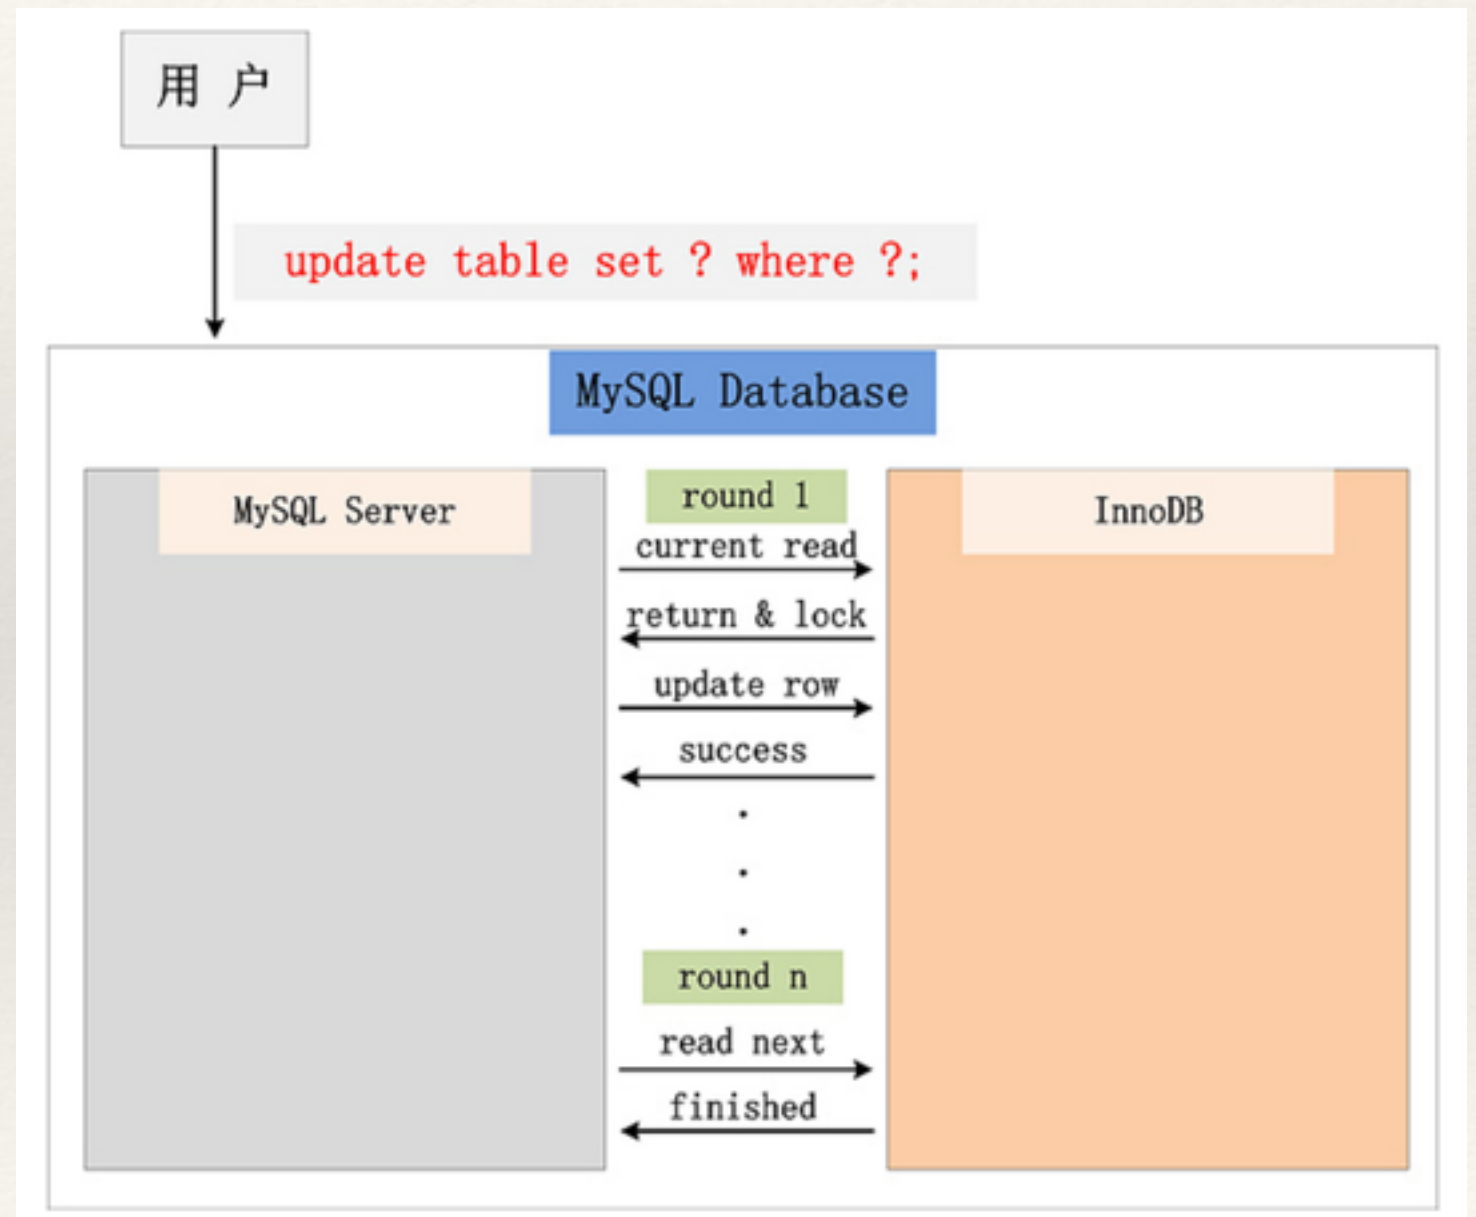
MVCC

- ❖ MVCC: Multi-Version Concurrency Control, compared Lock-Based Concurrency Control
- ❖ Advantage
 - ❖ read without lock
 - ❖ reading and writing is not conflict
- ❖ Isolation Level: RC、RR
- ❖ Read: Snapshot Read、Current Read

-
- 
-
- ❖ Snapshot Read: read history version、 unlock
 - ❖ select * from table where ...
 - ❖ Current Read: read latest version、 lock
 - ❖ select * from table where ... lock in share mode (S)
 - ❖ select * from table where ... for update (X)
 - ❖ insert into table values (...) (X)
 - ❖ update table set ... where ... (X)
 - ❖ delete from table where ... (X)


MVCC

- ❖ why update、delete、insert is current read?




Clustered Index

- ❖ Leaf: primary key value, transaction id, roll pointer, data
- ❖ not define, locates the first UNIQUE index where all the key columns are NOT NULL
- ❖ no suitable UNIQUE index, generates a 6-byte row ID



Secondary Index

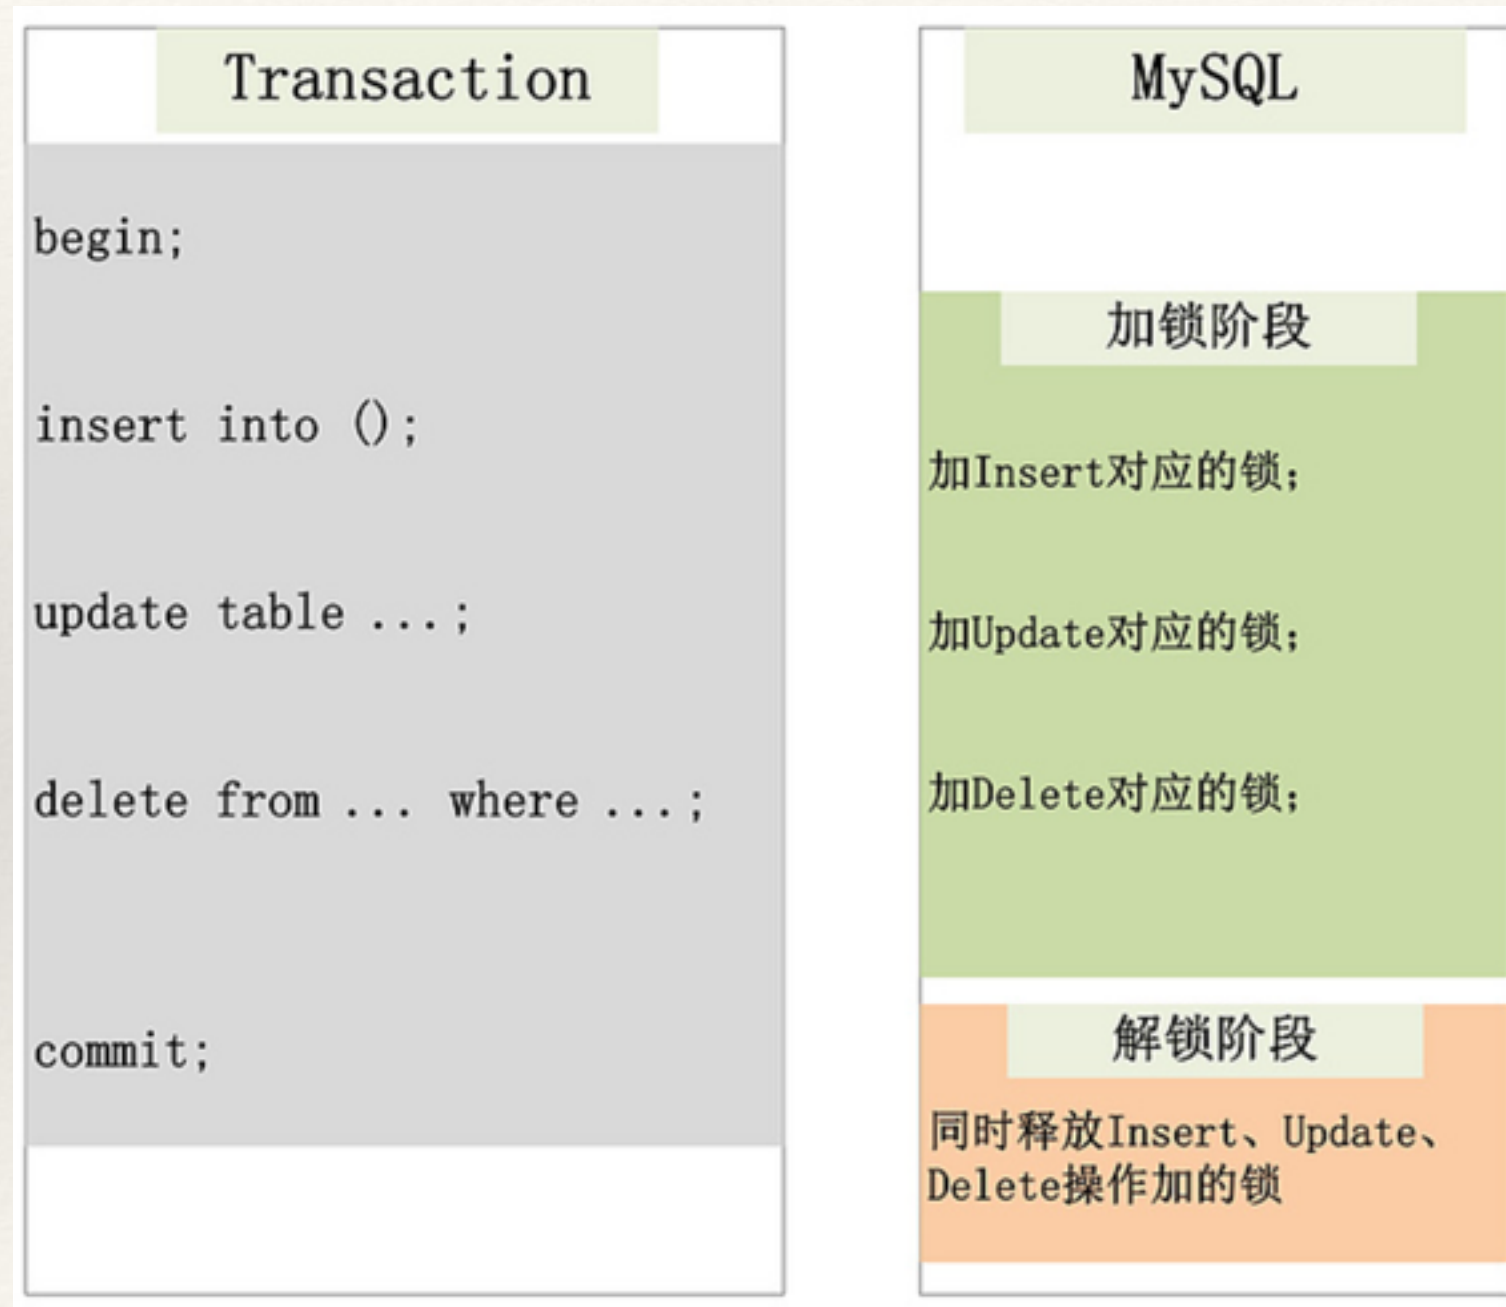
- ❖ Leaf: index Value、 primary key value
- ❖ Index Scan
- ❖ Secondary B+Tree find Primary Key
- ❖ Clustered B+Tree find data by Primary Key



Two-phase locking

- ❖ a concurrency control method that guarantees serializability
- ❖ Expanding phase
 - ❖ locks are acquired and no locks are released
- ❖ Shrinking phase
 - ❖ locks are released and no locks are acquired.

Two-phase locking



Isolation Level

- ❖ Read Uncommitted (RU)
- ❖ Read Committed (RC)
 - ❖ Snapshot Read: read latest version
 - ❖ Current Read: row lock, phantom read
- ❖ Repeatable Read (RR)
 - ❖ Snapshot Read: read transaction start version
 - ❖ Current Read: row lock + gap lock, solve phantom read
- ❖ Serializable
 - ❖ read operations are currently reading, read with read lock (S lock)
 - ❖ write with write lock (X lock)

Simple SQL Lock Realize Analyse

- ❖ SQL1: `select * from t1 where id = 10;`
- ❖ SQL2: `delete from t1 where id = 10;`
- ❖ Precondition
 - ❖ id is primary key、secondary index、unique index?
 - ❖ isolation level?
 - ❖ index scan? table scan?

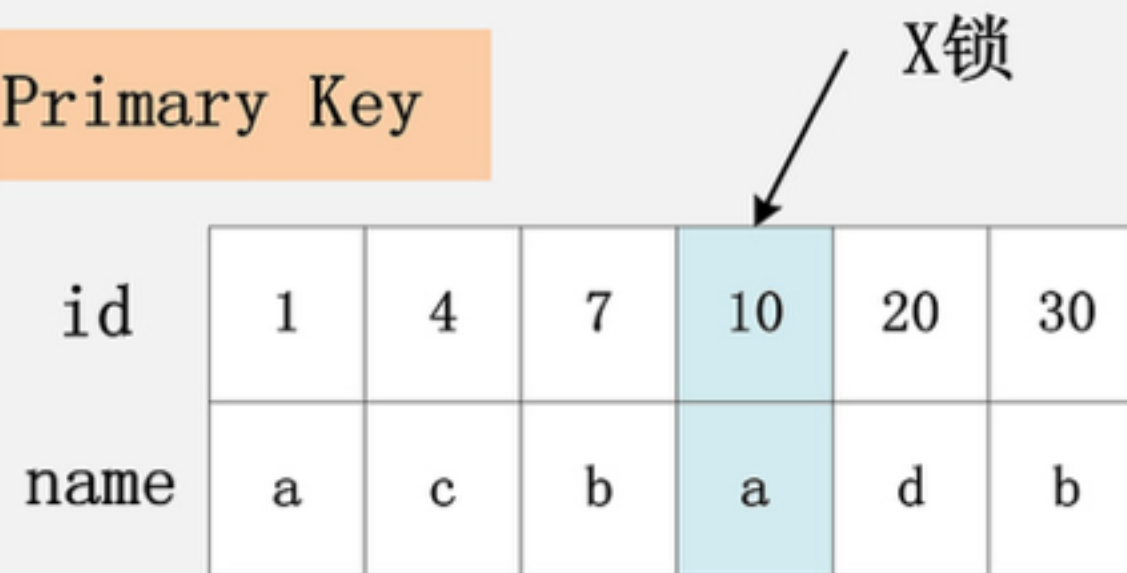
Simple SQL Lock Realize Analyse

- ❖ id PK + RC
- ❖ id Unique Index + RC
- ❖ id NonUnique Index + RC
- ❖ id Indexless + RC
- ❖ id PK + RR
- ❖ id Unique Index + RR
- ❖ id NonUnique Index + RR
- ❖ id Indexless + RR

PK + RC

Table: T1(id primary key, name)

Primary Key



The diagram illustrates an X lock (exclusive lock) being placed on the primary key value 10 in the table T1. An arrow points from the label 'X锁' to the cell containing '10' in the 'id' row. The cell containing 'a' in the 'name' row is also highlighted in light blue, indicating it is part of the locked record.

id	1	4	7	10	20	30
name	a	c	b	a	d	b

- ❖ PK id=10 record adds X lock

Unique Index + RC

Table: T1(name primary key, id unique key)

Unique Key (id)

id	1	2	3	5	6	10
name	f	zz	b	a	c	d

X锁

Primary Key

name	a	b	c	d	f	zz
id	5	3	6	10	1	2

X锁

- ❖ why PK is locked?
- ❖ update t1 set id = 100 where name = 'd'
- ❖ same record, delete and update serial execution

NonUnique Index + RC

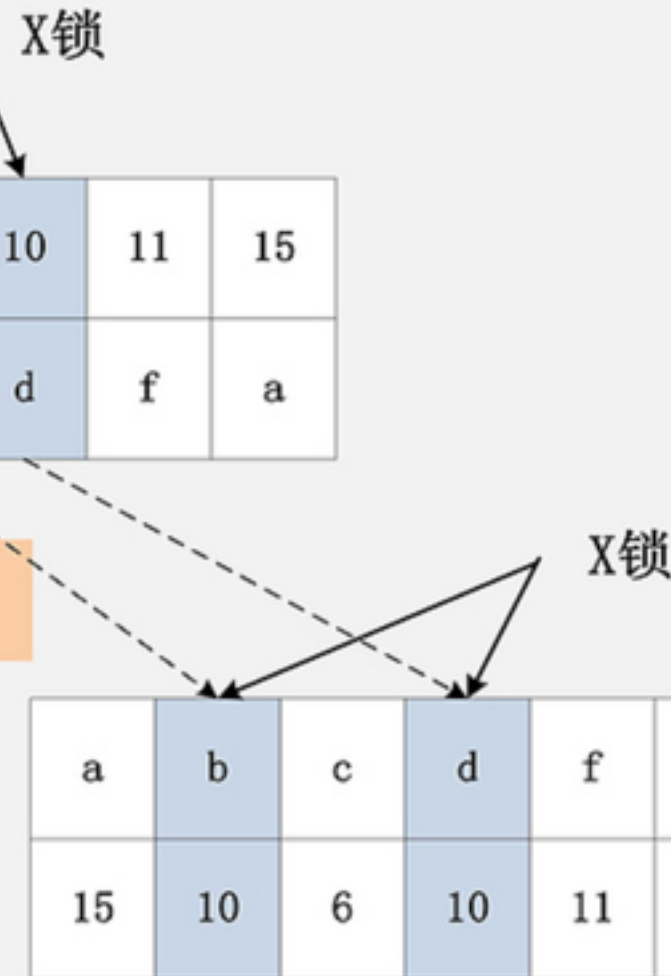
Table: T1(name primary key, id key)

Key (id)

id	2	6	10	10	11	15
name	zz	c	b	d	f	a

Primary Key

name	a	b	c	d	f	zz
id	15	10	6	10	11	2



No Index + RC

Table: T1(name primary key, id)

Primary Key

name	a	b	d	f	g	zz
id	5	3	10	2	10	9


- ❖ table scan by clustered index
- ❖ MySQL server will filter data, mysql optimize, semi-consistent read, go against the 2PL
- ❖ each record is locked and unlocked

PK + RR

Table: T1(id primary key, name)

Primary Key

X锁

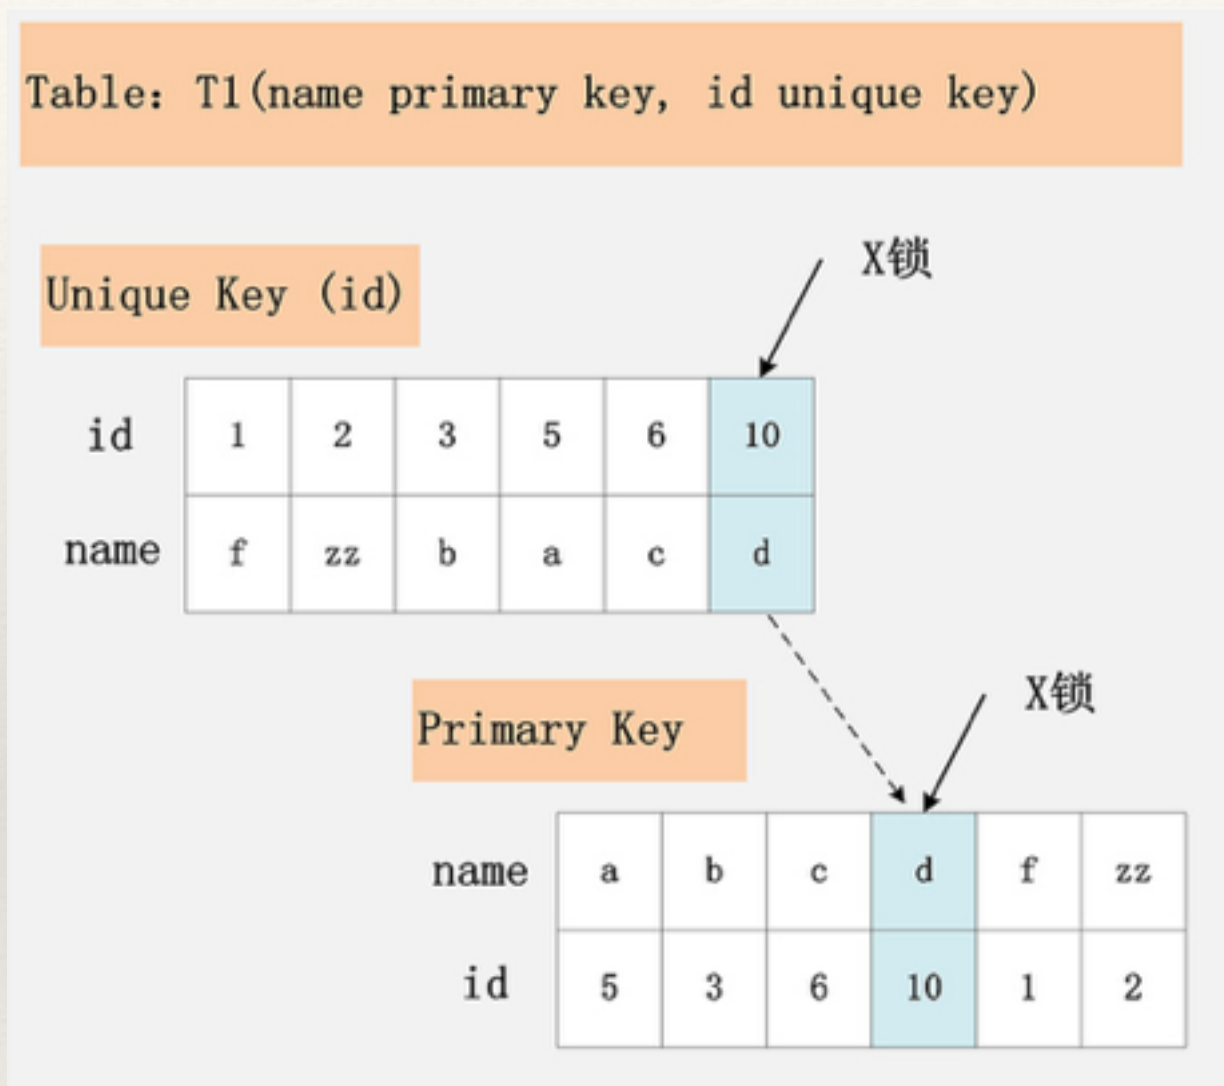


The diagram shows an arrow pointing from the text 'X锁' to the cell containing the value '10' in the 'id' column of the table. The entire row containing '10' and 'a' is highlighted in light blue.

id	1	4	7	10	20	30
name	a	c	b	a	d	b

- ❖ the same with PK + RC


Unique Index + RR



- ❖ the same with Unique Index + RC

NonUnique Index + RR

Table: T1(name primary key, id key)




NonUnique Index + RR

- ❖ what is phantom read?
- ❖ when needs the gap lock?
 - ❖ equality query
- ❖ considering :
 - ❖ `select * from t1 where id = 10 for update`
 - ❖ If not have records meeting the condition
 - ❖ `id PK + RR ?`
 - ❖ `id Unique Key + RR ?`

No Index + RR

Table: T1(name primary key, id)

Primary Key



- ❖ table scan
- ❖ semi-consistent read

Complicated SQL Lock Realize Analyse

Table: t1(id primary key, userid, blogid, pubtime, comment)

Index: idx_t1_pu(pubtime,userid)

idx_t1_pu

pubtime	1	3	5	10	20	100
userid	hdc	yyy	hdc	hdc	bbb	hdc
id	10	4	8	1	100	6

Primary Key

id	1	4	6	8	10	100
userid	hdc	yyy	hdc	hdc	hdc	bbb
blogid	a	b	c	d	e	f
pubtime	10	3	100	5	1	20
comment				good		

SQL: delete from t1 where pubtime > 1 and pubtime < 20 and userid = 'hdc' and comment is not NULL;

where conditions extraction


- ❖ index key
 - ❖ index first key, index last key
 - ❖ locates the index searching range
- ❖ index filter
 - ❖ each record filters
 - ❖ MySQL 5.6 brings in Index Condition Pushdown, before not distinguishes index filter and table filter,
- ❖ table filter
 - ❖ back to the clustered index to read the data

Complicated SQL Lock Realize Analyse

- ❖ Index key
 - ❖ `pubtime > 1 and puptime < 20`
 - ❖ `idx_t1_pu`
- ❖ Index Filter
 - ❖ `userid = 'hdc'`
 - ❖ `idx_t1_pu` filter
- ❖ Table Filter
 - ❖ `comment is not NULL`
 - ❖ clustered index filter

Complicated SQL Lock Realize Analyse

Table: t1(id primary key, userid, blogid, pubtime, comment)
Index: idx_t1_pu(pubtime,userid)



SQL: delete from t1 where pubtime > 1 and pubtime < 20 and userid = 'hdc' and comment is not NULL;

References

- ❖ <https://dev.mysql.com/doc>
- ❖ http://hedengcheng.com/?p=771#_Toc374698307
- ❖ <http://hedengcheng.com/?p=220>
- ❖ <http://hedengcheng.com/?p=577>



Thanks !